

CRISTOFF STAN

Senior Software Developer

Location: Tallinn, Harjumaa, Estonia | Email: cristoffstan@gmail.com
LinkedIn: <https://www.linkedin.com/in/cristoff-stan-866410383/> | Phone Number: +372 58619311

PROFESSIONAL SUMMARY

Innovative AI/ML Engineer and Full Stack Developer with 9+ years of hands-on experience architecting Python-driven backend solutions, LLM optimization workflows, and production-grade cloud-native systems. Skilled in Kubernetes, MLOps, and scalable infrastructure using AWS/GCP. Proven success in optimizing ML model performance, building fine-tuned data pipelines, and aligning AI systems to business goals. Adept at working in high-impact, remote-first environments using clean architecture and continuous delivery practices.

TECHNICAL SKILLS

- **Languages:** Python, Go, TypeScript, JavaScript, SQL
- **ML/DS:** LLM Tuning, Scikit-learn, Pandas, NumPy, MLOps, Model Monitoring
- **AI Ops & Infra:** Kubernetes, Docker, GitHub Actions, Prometheus, Grafana
- **Cloud Platforms:** AWS (EC2, S3, Lambda), GCP (Cloud Run, GKE), Vercel
- **Databases:** PostgreSQL, MySQL, MongoDB, Elasticsearch
- **Backend:** FastAPI, Flask, Django, gRPC
- **Frontend:** React, Next.js, Tailwind CSS
- **Dev Tools:** Git, JIRA, Postman, Pytest, Figma

WORK EXPERIENCE

Sensidev

Senior Software Engineer

September 2021 – December 2024

Timișoara, Romania

- Spearheaded AI initiatives including fine-tuning language models on ERP datasets using transfer learning techniques and orchestrated model pipelines via Dockerized FastAPI services.
- Engineered ML lifecycle tooling to monitor token usage, model latency, and retrieval accuracy for business intelligence dashboards integrating Elasticsearch.
- Developed cost-efficient AI inference services by introducing quantized versions of LLMs and using caching layers to optimize response latency for repetitive prompts.
- Integrated MLOps pipelines into existing CI/CD frameworks with GitHub Actions, enabling reproducible model training and versioned deployments to staging and production.
- Leveraged Kubernetes clusters to deploy scalable API endpoints, isolating memory-heavy models into dedicated pods with resource throttling.
- Led migration from rule-based systems to LLM-driven decision engines, aligning business workflows with NLP-based classification models.

- Integrated cloud-native logging (AWS CloudWatch & GCP Stackdriver) to capture real-time inference anomalies and trigger alerting for drift detection.

Syscom Digital

Software Engineer

October 2019 – August 2021

Bucharest, Romania

- Designed microservices in Python and FastAPI for real-time ML model execution on streaming marketing data, with outputs stored and indexed in PostgreSQL.
- Built and maintained analytics ETL jobs using Pandas for segment classification, anomaly flagging, and funnel performance tracking across large datasets.
- Implemented API gateways for model endpoints, enabling secure role-based access to predictions via OAuth2 and JWT.
- Participated in experiments benchmarking LLM response accuracy and contextual alignment on client-specific taxonomies.
- Used Docker Compose and Kubernetes manifests to simulate production load scenarios during model A/B testing and canary deployments.
- Assisted in building Grafana-based model performance dashboards by feeding prediction metrics via Prometheus exporters.
- Created alerting rules based on inference times and confidence thresholds to notify engineering teams of production incidents.

Soft Build

Software Developer

June 2016 – July 2019

Timișoara, Romania

- Developed Python tools to preprocess structured and unstructured HR data for sentiment analysis, feeding cleaned datasets into NLP models for internal insights.
- Refactored legacy code to adopt modular ML model definitions and transitioned monolithic training scripts into reusable, testable services.
- Automated deployment of model artifacts to internal endpoints using Linux shell scripts, cron jobs, and remote triggers.
- Collaborated on a POC for integrating ML scoring into Odoo ERP decisions, such as automated payroll anomaly detection and task prioritization.
- Created and maintained SQL data marts optimized for training set generation and slice-based validation.
- Improved feature engineering pipelines to support categorical encoding, imputation, and text vectorization using custom preprocessing layers.
- Designed job queuing system to orchestrate sequential model steps using Celery and Redis, managing compute cost and stability.

Decima Digital

Software Developer

November 2015 – April 2016

Tallinn, Estonia

- Supported foundational ML prototypes involving NLP on financial documents to classify transaction types with regex-enhanced tokenization strategies.
- Rewrote Python services to extract key entities from document streams and map them into normalized PostgreSQL schemas for supervised training.

- Participated in early-stage Kubernetes trials for deploying microservices and routing model responses through reverse proxies.
- Implemented in-house data validation pipelines for incoming model inputs and ensured consistent schema adherence prior to inference.
- Applied REST and RPC API design best practices for communicating prediction results back into Odoo-based dashboards.
- Used Jupyter notebooks to iterate quickly on model performance metrics and prepare decision-maker-ready visualizations.
- Performed data cleanup jobs to remove outliers from model input sources and maintain stable convergence.

Borna Technology

Software Developer Intern

September 2014 – August 2015

Tallinn, Estonia

- Prototyped FastAPI microservices to serve lightweight ML models for log classification and request metadata categorization.
- Assisted with proof-of-concept projects involving embedding-based similarity search on product names across ERP systems.
- Helped automate the retraining process for classification models using shell scripts, version tagging, and CSV data snapshots.
- Learned fundamentals of Kubernetes orchestration by shadowing DevOps engineers on container deployment practices.
- Supported model performance evaluation by logging precision/recall across multiple time slices using Matplotlib plots.
- Helped implement XML and JSON schema translators to align model inputs with different client formats.
- Built test harnesses for basic ML inference functions to validate logic prior to pipeline integration.

EDUCATION

Tallinn University

Bachelor's Degree in Computer Science,
2010 – 2014

CERTIFICATIONS

- ❖ **Deep Learning Specialization** – Andrew Ng (Coursera)
 - ❖ **LLM Engineering with LangChain & OpenAI** – DeepLearning.AI
 - ❖ **MLOps (Machine Learning Operations) Specialization** – Coursera
 - ❖ **Docker for Developers** – Udemy
 - ❖ **Certified Kubernetes Application Developer (CKAD)** – The Linux Foundation
 - ❖ **Google Cloud Professional Machine Learning Engineer** – Google
 - ❖ **Python for Data Science** – Coursera
 - ❖ **Odoo Technical Training** – OdooClass
-

PROJECTS

- ✓ **LLM Evaluation Pipeline** – Python-based benchmarking tool to test and compare LLM responses on enterprise knowledge prompts using structured metrics and feedback loops.
- ✓ **ML Cost Optimizer** – CLI and backend toolkit to simulate AI workload costs across cloud providers and recommend deployment strategies.
- ✓ **Model Drift Detector** – FastAPI microservice that scans real-time prediction data for statistical drift and triggers retraining alerts.