CRISTOFF STAN **Senior Software Developer**

Location: Tallinn, Harjumaa, Estonia | Email: cristoffstan@gmail.com
LinkedIn: https://www.linkedin.com/in/cristoff-stan-b53490362/ | Phone Number: +372 58619311

PROFESSIONAL SUMMARY

Innovative **Senior Software Engineer and Solution Architect** with 10+ years of experience designing, developing, and deploying **AI-driven applications, scalable architectures, and cloud-native systems**. Proven expertise in **Java, Node.js, and Python**, combined with hands-on experience in **AI/ML model integration, MLOps, and DevSecOps practices**. Adept at leading teams in **Agile environments**, gathering business requirements, and translating them into enterprise-grade solutions. Skilled in **AWS, GCP, CI/CD pipelines, and big data ecosystems**, with a strong track record of delivering high-impact digital solutions across international projects.

TECHNICAL SKILLS

- **Languages & Frameworks**: Java, Node.js, Python, TypeScript, JavaScript, C#, Angular, React, Spring Boot, Express.js, FastAPI, Flask
- **AI/ML & Data**: Machine Learning, Deep Learning, LLM Integration, RAG, TensorFlow, PyTorch, NLP, MLOps, Data Pipelines, Big Data Architectures
- **Cloud & DevOps**: AWS (EC2, S3, Lambda, Bedrock), GCP (Vertex AI, Pub/Sub, BigQuery), Azure, Kubernetes, Docker, Terraform, CloudFormation, Jenkins, GitLab CI/CD
- **Databases**: PostgreSQL, MySQL, SQL Server, MongoDB, Redis, Cassandra (SQL & NoSQL)
- **Tools & Practices**: JIRA, Git, Agile/Scrum, Microservices Architecture, Event-Driven Systems, Kafka, REST, GraphQL, API Gateway, Observability (Prometheus, Grafana)

WORK EXPERIENCE

Sensidev
Senior Software Engineer
September 2021 – December 2024
Timişoara, Romania

- Designed and implemented **AI-driven microservices in Python and Node.js**, integrating ML models for real-time data processing, anomaly detection, and predictive analytics.
- Led solution architecture for **enterprise applications** with microservices deployed on **AWS and GCP**, ensuring scalability, observability, and fault tolerance.
- Spearheaded adoption of **MLOps workflows**, automating training, deployment, and monitoring of ML models with pipelines using **CI/CD, Docker, and Kubernetes**.
- Partnered with product owners and business stakeholders to translate requirements into **solution blueprints and architecture diagrams** supporting mission-critical features.
- Introduced **DevSecOps best practices** with integrated security scanning, IaC (Terraform), and automated vulnerability checks across environments.

- Oversaw migration of legacy monolithic systems to **modern event-driven architectures** with **Kafka Streams** and distributed databases for high availability.
- Mentored and guided developers in applying **Agile engineering practices**, pair programming, and architectural reviews to improve delivery quality.
- Authored **comprehensive technical documentation** covering AI integration, APIs, and cloud architectures for long-term maintainability.

Syscom Digital
Software Engineer
October 2019 – August 2021
Bucharest, Romania

- Architected **cloud-based applications** combining Java, Python, and Node.js to support large-scale marketing automation systems with AI-powered personalization.
- Developed **NLP-based modules** for text classification, intent detection, and chatbot automation deployed via **AWS Lambda and GCP AI APIs**.
- Implemented **data ingestion pipelines** processing millions of records daily from external APIs into **PostgreSQL and MongoDB** with error resilience.
- Enhanced system reliability with **container orchestration in Kubernetes**, scaling workloads dynamically under peak user demand.
- Collaborated with cross-functional teams using **Agile/Scrum practices**, ensuring sprint deliverables included AI features aligned to business goals.
- Automated monitoring, logging, and alerting pipelines using **Prometheus, Grafana, and ELK Stack** for real-time observability.
- Introduced **GraphQL APIs** for optimized data access across multiple digital platforms, reducing query complexity and improving performance.
- Delivered technical specifications and architectural design patterns ensuring compliance with enterprise coding standards and future scalability.

Soft Build
Software Developer
June 2016 – July 2019
Timişoara, Romania

- Designed and built **AI-enabled recommendation systems** using Python and TensorFlow to enhance user personalization in e-commerce platforms.
- Implemented backend services in **Java and Node.js** connected to **SQL and NoSQL databases**, supporting millions of transactions monthly.
- Applied **cloud-native patterns** to deliver distributed services on **AWS EC2 and GCP App Engine**, improving global reach and system uptime.
- Migrated existing applications to **microservices architecture**, implementing APIs with REST and GraphQL for external integrations.
- Introduced **test automation pipelines** covering unit, integration, and end-to-end testing across multiple languages and frameworks.
- Optimized high-load services using **caching strategies in Redis** and performance profiling for low-latency response times.
- Documented system workflows, APIs, and deployment strategies to improve knowledge transfer and system supportability.
- Coordinated with stakeholders to align **AI project deliverables** with evolving business needs and regulatory requirements.

Decima Digital
Software Developer
November 2015 – April 2016
Tallinn, Estonia

- Developed and optimized enterprise **web and backend systems**, focusing on modular architecture to support future **AI/ML enhancements**.
- Designed **data access layers** that could scale for ML workloads, integrating SQL and NoSQL persistence strategies.
- Built **secure RESTful services** consumed by enterprise platforms with a focus on scalability, performance, and compliance.
- Partnered with data engineers to prepare structured and unstructured datasets for use in ML pipelines.
- Conducted performance tuning of database-heavy workloads, reducing query execution time by 40% on high-traffic systems.
- Designed **cloud deployment workflows** leveraging early AWS services to containerize applications for testing and staging environments.
- Collaborated with cross-functional teams to architect microservices with future ML-driven modules in mind.
- Produced detailed documentation and architectural guidelines for AI-readiness in customer platforms.

Borna Technology
Software Developer Intern
September 2014 – August 2015
Tallinn, Estonia

- Assisted in building early-stage **machine learning pipelines**, including feature extraction and preprocessing for text classification tasks.
- Implemented **search functionality with vector embeddings** to support semantic retrieval in prototype AI-driven systems.
- Developed internal **ETL workflows** for structured and unstructured data sources, preparing inputs for model training.
- Conducted experiments with **word embeddings and RNNs**, contributing to initial prototypes of natural language understanding systems.
- Automated **evaluation pipelines** with precision, recall, and F1 score metrics to validate model improvements.
- Built internal APIs for ML experimentation, integrating small-scale models into existing web applications.
- Optimized preprocessing steps with multi-threaded Python scripts, reducing dataset preparation time by 50%.
- Documented workflows and experimental results for use by senior engineers in productionizing ML solutions.

EDUCATION

Tallinn University

Bachelor’s Degree in Computer Science,
2010 – 2014

CERTIFICATIONS

- AWS Certified Machine Learning – Specialty
- Google Cloud Professional Machine Learning Engineer
- TensorFlow Developer Certificate
- DeepLearning.AI Large Language Models Specialization

PROJECTS

- ✓ **Enterprise RAG Assistant**: Built a multilingual Retrieval-Augmented Generation (RAG) pipeline with FAISS and LangChain, enabling Fortune 500 clients to deploy enterprise knowledge assistants under strict GDPR compliance.
- ✓ **LLM Optimization Framework**: Developed fine-tuning and distillation workflows for LLMs, applying quantisation and PEFT for low-latency deployment, reducing infrastructure costs by 40%.
- ✓ **Real-Time ML Monitoring**: Implemented an observability system tracking drift, accuracy, and inference latency across live ML services in AWS and GCP, enabling proactive retraining.