CRISTOFF STAN Senior Software Developer

Location: Tallinn, Harjumaa, Estonia | Email: cristoffstan@gmail.com
LinkedIn: https://www.linkedin.com/in/cristoff-stan-b53490362/ | Phone Number: +372 58619311

PROFESSIONAL SUMMARY

Seasoned MLOps Engineer and Senior Software Developer with over 9 years of experience designing, deploying, and automating scalable machine learning infrastructure. Proven success in managing the complete ML lifecycle—from data workflows and model training pipelines to CI/CD orchestration and real-time production monitoring. Highly proficient in Kubernetes, Docker, Terraform, and multi-cloud architectures (GCP, AWS). Adept at enabling seamless collaboration between data science and engineering teams while implementing best practices in MLOps, observability, and reproducibility.

TECHNICAL SKILLS

- **Languages:** Python, Bash, SQL, Go (basic)
- **ML Tooling:** MLflow, TFX, Weights & Biases, Metaflow
- **Orchestration:** Airflow, Prefect, Kubeflow Pipelines
- **Cloud & Infra:** GCP, AWS, Azure, Kubernetes (EKS, GKE, AKS), Docker, Terraform, Helm
- **Data & Messaging:** Kafka, Spark, Snowflake, dbt, PostgreSQL, Redis
- **Monitoring & Logging:** Prometheus, Grafana, ELK Stack, Datadog
- **CI/CD & DevOps:** GitHub Actions, GitLab CI, Jenkins, ArgoCD, ML model versioning
- **Other:** Feature stores, model registries, vector databases, Responsible AI, LLM workflows

---

WORK EXPERIENCE

Sensidev
Senior Software Engineer
September 2021 – December 2024
Timişoara, Romania

- Designed and deployed robust CI/CD pipelines for ML models using GitHub Actions, Terraform, and Kubernetes to support large-scale product deployments across GCP and AWS.
- Built custom Airflow DAGs to orchestrate data pipelines and ML model retraining triggered by performance drift and A/B test outcomes.
- Led containerization and deployment of ML microservices using Docker and Helm charts, achieving 40% improvement in model delivery time.
- Integrated MLflow for experiment tracking, model registry, and lifecycle automation across development and production environments.
- Collaborated with data scientists to enable reproducible ML workflows and implemented feature store integration for dynamic input pipelines.
- Established monitoring with Prometheus and Grafana to visualize inference latency, model accuracy, and system anomalies.

- Implemented infrastructure-as-code strategy using Terraform modules to standardize cloud provisioning across teams and regions.
- Championed the adoption of Responsible AI governance frameworks and automated model audit logs for regulatory compliance.

Syscom Digital
Software Engineer
October 2019 – August 2021
Bucharest, Romania

- Developed and maintained distributed ML pipelines using Kubeflow and Prefect to enable real-time prediction systems and batch scoring jobs.
- Engineered containerized data ingestion and transformation jobs that ingested over 5TB/day into Snowflake, powering downstream analytics.
- Created automated CI/CD workflows for training and deploying NLP models via Docker and Jenkins pipelines.
- Integrated Weights & Biases for hyperparameter tuning and experimentation visualization to accelerate iteration cycles.
- Managed cloud-native deployments in Kubernetes (EKS) and orchestrated secure model rollout procedures across staging and prod clusters.
- Led data versioning strategy combining dbt with GitOps practices for reproducible model input lineage.
- Developed observability layers using Datadog and ELK to provide SLA tracking and real-time alerts on model accuracy and availability.
- Collaborated cross-functionally with backend and DevOps teams to align ML infrastructure with system-wide engineering practices.

Soft Build
Software Developer
June 2016 – July 2019
Timişoara, Romania

- Designed internal tooling in Python for model performance benchmarking and automated test harness generation.
- Implemented data validation pipelines using Python and Great Expectations to ensure input data integrity before model training.
- Developed secure APIs for ML inference with FastAPI and integrated authentication and monitoring hooks for auditing usage.
- Participated in migrating legacy ML services to container-based deployments in Docker and Kubernetes.
- Built scheduled retraining workflows with model drift detection logic embedded via Spark and scikit-learn pipelines.
- Maintained real-time prediction endpoints using RESTful architecture and optimized response time through async caching techniques.
- Conducted performance tuning on Kubernetes clusters, reducing pod scaling lag by optimizing resource quotas and autoscalers.
- Wrote internal documentation and onboarding guides to train new engineers on ML infrastructure best practices.

Decima Digital
Software Developer
November 2015 – April 2016
Tallinn, Estonia

- Built ETL scripts in Python and Bash for data preprocessing and pipeline testing.
- Assisted in model deployment on AWS EC2 instances and created custom scripts for model restarts and logging.
- Designed schema migration plans and created Terraform modules for small-scale infrastructure setups.
- Monitored cloud instances using CloudWatch and implemented basic log parsing scripts for model output review.
- Participated in sprint planning to scope MLOps automation features with a team of engineers and data analysts.
- Created user-friendly CLI wrappers for pipeline operations to simplify ML workflows for non-technical stakeholders.
- Set up containerization for Python inference services and documented Dockerfile best practices for shared use.
- Conducted exploratory data analysis and presented findings on model training bottlenecks to team leads.

Borna Technology
Software Developer Intern
September 2014 – August 2015
Tallinn, Estonia

- Supported ML engineers by automating dataset versioning and implementing CLI-based deployment scripts.
- Wrote Bash and Python scripts to manage training jobs on cloud instances, ensuring GPU resource optimization.
- Documented ML pipeline configuration and performance logs across several experimental projects.
- Participated in data labeling and validation efforts to support initial supervised learning initiatives.
- Automated CSV validation and formatting scripts used in pre-training processes.
- Collaborated on pipeline stress tests and documented model latency under different data loads.
- Attended internal workshops on containerization and ML lifecycle management.
- Delivered final internship presentation on infrastructure optimization for ML training experiments.

EDUCATION

Tallinn University

Bachelor’s Degree in Computer Science,
2010 – 2014

CERTIFICATIONS

- MLOps Specialization – DeepLearning.AI / Coursera
- Google Cloud Professional Machine Learning Engineer
- Kubernetes for Machine Learning – Udacity
- AWS Certified Solutions Architect – Associate

PROJECTS

- ✓ **ML Infra at Scale** – Designed a multi-cloud (GCP + AWS) MLOps infrastructure with Terraform, Helm, and MLflow. Integrated retraining triggers, model registries, and Airflow orchestration.
- ✓ **Real-Time NLP Monitoring** – Built and deployed an NLP sentiment analysis API with real-time Prometheus monitoring and Grafana dashboards to track usage and model drift.
- ✓ **CI/CD for ML Microservices** – Implemented GitHub Actions-based CI/CD pipelines for ML model training and deployment, reducing time-to-production by 50% and boosting reproducibility.