

CRISTOFF STAN

Senior Software Engineer

Location: Tallinn, Harjumaa, Estonia | Email: cristoffstan@gmail.com
LinkedIn: <https://www.linkedin.com/in/cristoff-stan-b53490362/> | Phone Number: +372 58619311

PROFESSIONAL SUMMARY

Results-driven Senior Data Engineer with over 9 years of experience designing, developing, and deploying large-scale data solutions. Expert in SQL, Python, API integrations, and modern data platforms including Snowflake, BigQuery, and dbt. Proven track record of leading cross-functional teams, mentoring junior engineers, and delivering impactful insights for business decision-making. Experienced with MLOps, GenAI techniques, and building AI-augmented pipelines that enable intelligent data-driven workflows.

TECHNICAL SKILLS

- Programming & Data: Python, SQL, Pandas, NumPy, dbt, Apache Airflow, Jupyter
- Cloud & Warehousing: Snowflake, BigQuery, Fabric, AWS (Redshift, Lambda), Azure, GCP
- Data Visualization: Power BI, Looker, Metabase, Steep
- ML & AI: Scikit-learn, OpenAI APIs, Hugging Face, LangChain, MLflow, Vertex AI, GenAI architecture
- DevOps & APIs: REST APIs, Git, Docker, CI/CD, FastAPI, Postman
- Other Tools: Fivetran, Kafka, Terraform, Jenkins, Notion, Jira

WORK EXPERIENCE

Senior Software Engineer

Sensidev

September 2021 – December 2024

Timișoara, Romania (Remote)

- Architected and deployed end-to-end data pipelines in Snowflake and BigQuery, integrating dbt models to transform raw data into analytics-ready formats for international clients.
- Built and maintained robust ELT frameworks using Python and Airflow, improving data latency by 40% across batch processing workflows.
- Led development of internal tools for automating data quality checks, schema validation, and anomaly detection using Python and Metabase dashboards.
- Designed and implemented scalable APIs to expose analytical data models for client applications, ensuring consistent integration across BI and ML platforms.
- Collaborated with data scientists to deploy GenAI-enabled modules that provided contextual data summarizations for stakeholders in finance and retail.
- Managed a team of 4 data engineers, conducting regular mentorship and peer-review sessions to promote clean code, CI/CD practices, and dbt documentation.

- Spearheaded a migration project from on-prem PostgreSQL to BigQuery, achieving a 75% reduction in query execution time for reporting workloads.
- Partnered with executive leadership to scope new data consulting engagements, building client-specific roadmaps and prototypes using Looker and Power BI.

Software Engineer

Syscom Digital

October 2019 – August 2021

Bucharest, Romania (Remote)

- Developed Python-based backend services to process, analyze, and transform campaign engagement data from multiple client sources.
- Implemented SQL-based analytics dashboards in Power BI and Metabase to track campaign performance, A/B test results, and segmentation efficiency.
- Designed RESTful APIs and integrated with third-party systems, including CRMs and ad platforms, ensuring seamless data exchange and automation of marketing workflows.
- Created dbt models to standardize and modularize logic for KPIs, streamlining dashboard maintenance and client-specific reporting.
- Built custom logging and alerting mechanisms for data pipeline failures using Python, Slack API, and Airflow, achieving 90% reduction in downtime.
- Introduced data versioning strategies with Git and CI pipelines, enabling faster collaboration and reproducibility across the analytics team.
- Supported MLOps pipelines for predictive models used in user segmentation and recommendation, using MLflow for model lifecycle management.
- Conducted technical training sessions on SQL optimization, Power BI, and dbt best practices for internal analysts and junior engineers.

Software Developer

Soft Build

June 2016 – July 2019

Timișoara, Romania (Remote)

- Engineered internal data platforms using Python and PostgreSQL, supporting a user analytics engine for SaaS product usage insights.
- Implemented client-specific dashboards in Power BI and Looker, embedding dynamic filters and real-time refresh logic using REST API calls.
- Designed modular extract-transform-load (ETL) routines that ingested data from 12+ external APIs, normalizing inputs using dbt staging models.
- Participated in the migration of legacy analytics systems to cloud-native infrastructure on AWS with Redshift, reducing hosting costs by 60%.
- Built role-based access controls across multiple analytics tools to enforce GDPR-compliant data access policies for client organizations.
- Developed a Python SDK wrapper for internal REST APIs, improving developer productivity and decreasing onboarding time for new hires.
- Created test harnesses and data mocks to improve integration test coverage across analytical services, contributing to a 35% drop in bugs.
- Contributed to design and documentation of analytics SLAs, leading to more structured hand-offs with external business teams and clients.

Software Developer
Decima Digital
November 2015 – April 2016
Tallinn, Estonia

- Built backend services using Python and Flask to aggregate client sales data into a centralized analytics repository.
- Integrated REST APIs for 6 e-commerce platforms, standardizing product and order data for unified business reporting.
- Developed SQL views and indexes to improve performance of frequently accessed dashboards, reducing query load by 50%.
- Produced daily data quality reports delivered via email and Slack integrations, helping business analysts take timely action.
- Collaborated with frontend developers to design APIs for interactive data visualizations in admin panels.
- Participated in agile sprint planning and client demos, ensuring delivery alignment with business needs and expectations.
- Introduced Python scripts to clean legacy CSV datasets, handling edge cases and data inconsistency issues effectively.
- Maintained GitHub repositories, resolving issues and performing pull request reviews in a collaborative team setting.

Software Developer Intern
Borna Technology
September 2014 – August 2015
Tallinn, Estonia (On-Site)

- Created scripts for transforming raw sensor data from IoT devices into structured formats for dashboard reporting.
- Worked on ingestion pipelines using Python and SQLite, preparing data for analytical use in pilot projects.
- Built prototype dashboards in Excel and Power BI to visualize metrics such as uptime, temperature fluctuation, and energy use.
- Conducted benchmark comparisons for cloud storage providers to determine optimal options for client telemetry.
- Implemented unit tests and debugging logs to ensure robustness in data formatting and transmission logic.
- Assisted senior engineers in deploying APIs and running local test environments for data simulation.
- Researched best practices in IoT analytics and presented findings during internal workshops.
- Contributed to technical documentation and onboarding materials for incoming interns and junior staff.

EDUCATION

Tallinn University

Bachelor's Degree in Computer Science,
2010 – 2014

CERTIFICATIONS

- dbt Fundamentals Certification (dbt Labs)
 - Google Cloud Certified: Professional Data Engineer
 - Microsoft Certified: Azure Data Engineer Associate
 - Data Science & Machine Learning Bootcamp (Udemy)
 - OpenAI GPT-4 Developer Specialization (OpenAI Academy)
-

PROJECTS

- ✓ **AI Workflow Augmentation:** Led development of a GenAI-based knowledge assistant using OpenAI APIs to surface insights from unstructured business documentation. Integrated with Slack and Power BI.
- ✓ **Marketing Data Hub:** Built a centralized warehouse integrating Google Ads, Meta, and HubSpot APIs via Fivetran and dbt, used by 15 clients for real-time campaign optimization.
- ✓ **Smart KPI Framework:** Designed modular KPI modeling framework using dbt and Metabase that auto-adjusts based on business vertical and client segmentation.
- ✓ **ML-Powered Retention Engine:** Collaborated with data scientists to productionize churn prediction models using Scikit-learn and MLflow, saving \$200K/year in customer loss.
- ✓ **GDPR Compliance Engine:** Developed APIs to automatically flag and redact sensitive PII across data lakes using Python and Regex detection rules.