

PROFESSIONAL SUMMARY

Senior **Data Warehouse engineer** with 12+ years across SaaS, property technology, nonprofit platforms, civic technology, and IT services, specializing in **Python 3.10-3.12, SQL, Airflow 2.x, dbt, PostgreSQL, Snowflake, AWS** for secure, scalable platforms; delivers **architecture** ownership, modernization, **performance** tuning, **test automation**, **CI/CD** discipline, and **production reliability** that improves customer workflows, **release quality**, and engineering velocity.

TECH SKILLS

- Data Engineering: **Python 3.10-3.12, SQL, Airflow 2.x, dbt, PostgreSQL, Snowflake, AWS**, batch and streaming pipelines, data modeling, dimensional design, ELT/ETL orchestration, data quality rules
- Databases & Warehouses: **PostgreSQL 13-16, MySQL 8, SQL Server 2019-2022, Snowflake**, BigQuery, Redshift, Delta Lake, query optimization
- Cloud & Orchestration: **AWS, Azure, GCP, Airflow 2.x, dbt**, Spark 3.x, Databricks, **Kafka**, Glue, Data Factory, **CI/CD** for data pipelines
- BI & Analytics: Power BI, Tableau, semantic layers, KPI governance, dashboard **performance**, metric definitions, stakeholder reporting workflows
- Senior Practices: Data contracts, lineage, **observability**, cost control, privacy controls, incident recovery, schema evolution, cross-functional delivery

PROFESSIONAL EXPERIENCE

Staff Software Engineer

Relatio

Jun 2022 - Apr 2026

Las Vegas, NV

- Led data platform improvements for a SaaS workflow and analytics product using **Python 3.10-3.12, SQL, Airflow 2.x, dbt, PostgreSQL, Snowflake, AWS**, improving reporting freshness by **28%** through better orchestration, incremental loads, and data quality checks.
- Designed trusted data models, transformation layers, and reconciliation workflows that reduced dashboard mismatch defects by **19%** across customer reporting modules.
- Modernized legacy batch jobs into observable pipelines with retry handling, lineage notes, schema checks, and failure alerts for high-volume operational datasets.
- Resolved production issues involving slow warehouse queries, duplicate records, missing partitions, late-arriving events, and inconsistent customer-level aggregation rules.
- Implemented automated validation for source extracts, transformations, and BI-ready tables so data releases could move through **CI/CD** with stronger confidence.
- Optimized **SQL**, indexing, clustering, and materialized views for workflow analytics, reducing heavy report wait times during peak customer usage.
- Partnered with product and finance teams to define metric logic, data contracts, retention rules, and privacy controls for enterprise reporting features.
- Created **observability** dashboards for pipeline duration, freshness, failure rates, row-count anomalies, and warehouse cost so incidents were easier to diagnose.
- Mentored engineers on data modeling, testable transformations, pipeline reliability, stakeholder communication, and responsible handling of sensitive customer data.

Senior Software Engineer

AppFolio

Feb 2018 - May 2022

Goleta, CA

- Built property technology data pipelines using earlier versions of **Python 3.10-3.12, SQL, Airflow 2.x, dbt, PostgreSQL, Snowflake, AWS**, improving rental reporting processing speed by **24%** through better **SQL** models and scheduled workflow design.
- Implemented tenant communication, payment, lease, and document datasets with validation, transaction logging, reconciliation checks, and support-friendly error reporting.
- Migrated older reporting jobs from fragile scripts into reusable pipeline components with stronger testing and clearer ownership between product and data teams.
- Fixed defects involving duplicate payment events, delayed notification records, inconsistent lease statuses, and timeout errors by adding idempotency and data constraints.
- Optimized warehouse queries, source indexes, and cached lookup datasets used by high-read reporting pages and internal customer support workflows.
- Improved data quality coverage with unit, integration, and reconciliation tests tied to CI pipelines for high-volume property management features.
- Collaborated with QA and product teams to reproduce hard-to-detect data defects, then delivered targeted fixes without introducing broader regression risk.
- Mentored engineers on **SQL performance**, pipeline troubleshooting, data modeling, and pragmatic debugging during feature delivery and production support.

Software Engineer

Blackbaud

Jul 2016 - Jan 2018

Charleston, SC

- Developed nonprofit data workflows for fundraising, donor management, imports, reporting, and activity history using relational models and message-based processing.
- Improved donor search and reporting workflows by tuning database queries and transformation logic, reducing common reporting delays by **21%** for large datasets.
- Implemented secure data import services with validation rules, duplicate detection, row-level error reporting, and rollback-safe processing for support teams.
- Resolved bugs related to malformed CSV imports, timezone conversion issues, partial transaction failures, and mismatched campaign totals by adding stricter parsing.
- Built reusable data components for campaign records, donor profiles, activity history, and aggregate reporting with consistent backend behavior.

- Added automated tests around donation, reporting, and import workflows, improving confidence before releases and reducing regression defects.
- Worked with DevOps to improve deployment checks, environment configuration, and logging so data-related production issues could be identified faster.

Software Developer

Granicus

May 2014 - Jun 2016

Dallas, TX

- Built civic technology reporting and data processing features for agenda management, public records, document metadata, and subscription workflows.
- Improved public meeting content processing reliability through better job scheduling, retry handling, metadata validation, and backend data checks.
- Implemented reporting endpoints and data extraction workflows while keeping permission checks and customer configuration consistent across internal modules.
- Fixed recurring production issues involving failed document imports, missing metadata, slow search responses, and stale reporting rows by improving validation and indexing.
- Created integration tests for data workflows that handled public records, meeting documents, and notification delivery before production releases.
- Worked with support teams to investigate customer incidents, reproduce data errors, and deliver safe patches without disrupting active government portals.

Junior Software Developer

Vology

Jun 2013 - Apr 2014

Gainesville, FL

- Supported internal reporting applications using **SQL**, **Python** scripts, **MySQL**, shell automation, and Linux tools for operational data maintenance.
- Built small data utilities for ticket routing, inventory updates, and operational reporting that reduced repetitive manual work for internal teams.
- Fixed bugs in legacy scripts involving bad input handling, missing null checks, inconsistent database updates, and incomplete report totals by adding validation.
- Assisted senior engineers with database scripts, data cleanup tasks, and deployment preparation while learning production debugging and software delivery practices.
- Wrote unit tests, technical notes, and troubleshooting guides for data utilities, helping the team maintain older systems with fewer support escalations.

SELECTED PROJECTS

- Operational Data Platform: Built governed pipelines with **Python 3.10-3.12**, **SQL**, **Airflow 2.x**, **dbt**, **PostgreSQL**, **Snowflake**, **AWS**, data quality checks, orchestration, and dimensional models that improved reporting trust and reduced manual reconciliation for business users.
- Analytics **Performance** Modernization: Tuned warehouse queries, semantic metrics, dashboards, and scheduled jobs with **Python 3.10-3.12**, **SQL**, **Airflow 2.x**, **dbt**, **PostgreSQL**, **Snowflake**, **AWS**, enabling faster executive reporting and more reliable **customer-facing** analytics delivery.

EDUCATION

Arkansas State University - Bachelor's Degree in Computer Science, Sep 2009 - May 2013